

Meeting 6 SP23
Intro to Pipelining

@crawfish



Outline

Motivation

Pipelining

What else?



Section 1

Motivation



Goals of Processor Design

- Low latency
- High throughput
- Small size
- Low energy requirements
- etc.



Latency vs. Throughput

- **Latency** refers to the time needed to completely a single task



Latency vs. Throughput

- **Latency** refers to the time needed to completely a single task
- **Throughput** refers to the average number of tasks we can complete per unit time



Latency vs. Throughput

- **Latency** refers to the time needed to completely a single task
- **Throughput** refers to the average number of tasks we can complete per unit time
 - ▶ Most relevant for program run-time



Latency vs. Throughput

- **Latency** refers to the time needed to completely a single task
- **Throughput** refers to the average number of tasks we can complete per unit time
 - ▶ Most relevant for program run-time
- Pipelined processor designs optimize for throughput
 - ▶ Extra logic required actually causes a slight *increase* in latency

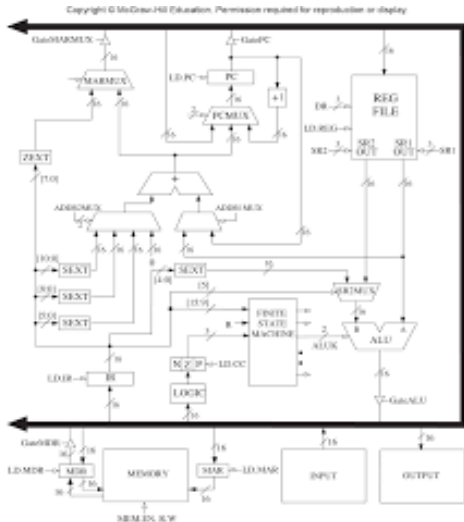


Single Cycle Processors

A **single cycle processor** is one in which an individual instruction takes one clock cycle to execute*.



LC3 Single-Cycle Datapath



Why is this Slow?

- Only one instruction is being processed per clock cycle



Why is this Slow?

- Only one instruction is being processed per clock cycle
- Clock cycles are long



Section 2

Pipelining



Multi-Cycle Processors

A **multi-cycle processor** takes multiple clock cycles to execute a single instruction



Multi-Cycle Processors

A **multi-cycle processor** takes multiple clock cycles to execute a single instruction

- Pipelined processors are a type of multi-cycle processor in which the CPU can have multiple in-progress instructions at once



Multi-Cycle Processors

A **multi-cycle processor** takes multiple clock cycles to execute a single instruction

- Pipelined processors are a type of multi-cycle processor in which the CPU can have multiple in-progress instructions at once
- Instruction handling is partitioned into 'units'




Multi-Cycle Processors

A **multi-cycle processor** takes multiple clock cycles to execute a single instruction

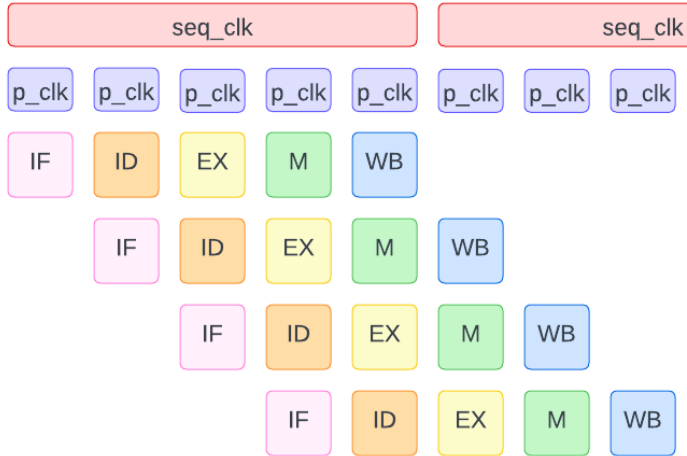
- Pipelined processors are a type of multi-cycle processor in which the CPU can have multiple in-progress instructions at once
- Instruction handling is partitioned into 'units'
- Instruction dispatch is staggered such that each unit is always working on exactly 1 instruction



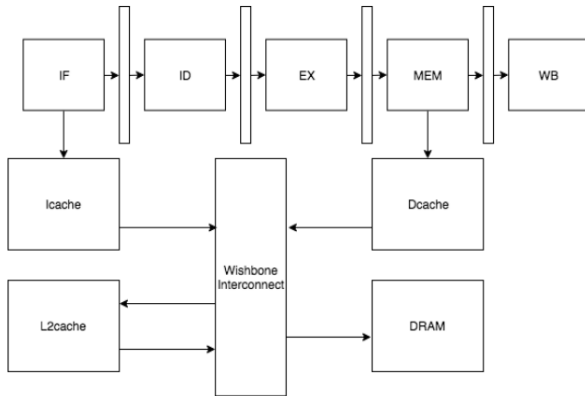
Sequential Execution



Pipelining



LC3 Pipelined Architecture



Instruction Fetch

```
module InstructionFetch(  
    input clk,  
    input [15:0] pc,  
    input stall,  
    output [15:0] instr_out,  
  
    // One way read-only memory itf  
    mem_itf.read_only_o mem_itf  
);
```



Instruction Decode

```
module InstructionDecode(  
    input [15:0] instr,  
  
    // From instr 'microcode', produce ctrl signals  
    control_signals_itf.control_signals cs_itf  
);
```



Execution

```
module ExecutionUnit(  
    // reg_a, reg_b are read from register file  
    input logic [15:0] reg_a,  
    input logic [15:0] reg_b,  
    input logic [1:0] opcode,  
  
    // c is forwarded for later writeback  
    output logic [15:0] c  
);
```



Memory Ops

- Interactions with memory occur here
 - ▶ Load / store instructions



Memory Ops

- Interactions with memory occur here
 - ▶ Load / store instructions
- Memory operations are not single cycle (even w/ caching)



Memory Ops

- Interactions with memory occur here
 - ▶ Load / store instructions
- Memory operations are not single cycle (even w/ caching)
- Less regular access patterns than instruction fetch



Memory Ops

- Interactions with memory occur here
 - ▶ Load / store instructions
- Memory operations are not single cycle (even w/ caching)
- Less regular access patterns than instruction fetch
- Memory operations are typically less frequent than arithmetic operations



Memory Ops

- Interactions with memory occur here
 - ▶ Load / store instructions
- Memory operations are not single cycle (even w/ caching)
- Less regular access patterns than instruction fetch
- Memory operations are typically less frequent than arithmetic operations
 - ▶ Reducing the need to interact with memory is a huge area of optimization



Memory Ops

- Interactions with memory occur here
 - ▶ Load / store instructions
- Memory operations are not single cycle (even w/ caching)
- Less regular access patterns than instruction fetch
- Memory operations are typically less frequent than arithmetic operations
 - ▶ Reducing the need to interact with memory is a huge area of optimization
- How do we deal w/ long memory operations in our pipeline?



Memory Ops

- Interactions with memory occur here
 - ▶ Load / store instructions
- Memory operations are not single cycle (even w/ caching)
- Less regular access patterns than instruction fetch
- Memory operations are typically less frequent than arithmetic operations
 - ▶ Reducing the need to interact with memory is a huge area of optimization
- How do we deal w/ long memory operations in our pipeline?
 - ▶ Stall :(



Write Back

- Instruction results are 'committed' to register file
- Input comes from either execution stage or register file depending on control signals



What are some problems with this design?




```
; R4 = 0x1, R5 = 0x2
```

```
ADD R4, R5, R0
```

```
ADD R4, R4, R4
```

```
NOP
```

```
NOP
```

```
NOP ; R4 = 0x6?
```



Hazards

- **RAW:** Read After Write
 - ▶ If we 'read' data from the register file (or memory) too soon after it is written, we may encounter a race condition where the returned data is stale
 - ▶ Write has not yet been 'committed'



Hazards

- **RAW:** Read After Write
 - ▶ If we 'read' data from the register file (or memory) too soon after it is written, we may encounter a race condition where the returned data is stale
 - ▶ Write has not yet been 'committed'
- How do we fix this?



Hazards

- **RAW:** Read After Write
 - ▶ If we 'read' data from the register file (or memory) too soon after it is written, we may encounter a race condition where the returned data is stale
 - ▶ Write has not yet been 'committed'
- How do we fix this?
 - ▶ Wait



Hazards

- **RAW:** Read After Write
 - ▶ If we 'read' data from the register file (or memory) too soon after it is written, we may encounter a race condition where the returned data is stale
 - ▶ Write has not yet been 'committed'
- How do we fix this?
 - ▶ Wait
 - ▶ Instruction forwarding
- Other hazard types exist in more complicated pipeline structures



Branch and Jump Instructions

- Branch / Jump instructions modify the program counter in an unexpected* way



Branch and Jump Instructions

- Branch / Jump instructions modify the program counter in an unexpected* way
- These changes take effect during the 'write-back' phase



Branch and Jump Instructions

- Branch / Jump instructions modify the program counter in an unexpected* way
- These changes take effect during the 'write-back' phase
 - ▶ Implemented as an addition or replacement of the program counter register



Branch and Jump Instructions

- Branch / Jump instructions modify the program counter in an unexpected* way
- These changes take effect during the 'write-back' phase
 - ▶ Implemented as an addition or replacement of the program counter register
- What about instructions that populate the pipeline in the meantime?
 - ▶ Not useful instructions
 - ▶ If they modify state (registers, memory) they will cause incorrect future results upon completion



Branch and Jump Instructions

- Stall the pipeline (Wait)



Branch and Jump Instructions

- Stall the pipeline (Wait)
- Branch prediction (Guess)
 - ▶ If we guessed wrong, don't commit anything



Branch and Jump Instructions

- Stall the pipeline (Wait)
- Branch prediction (Guess)
 - ▶ If we guessed wrong, don't commit anything
 - ▶ Lots of research has been done to optimize branch prediction
 - ▶ Code is very repetitive
 - ▶ If we guess, we can be right fairly often



Golden Rule of CPU Design: Hide Latency to Increase Throughput



The end. (Now go do some 233 practice problems)

